

Main Points



- Process concept 
 - A process is the OS abstraction for executing a program with limited privileges
- Dual-mode operation: user vs. kernel 
 - Kernel-mode: execute with complete privileges
 - User-mode: execute with fewer privileges
- Safe control transfer
 - How do we switch from one mode to the other?

Mode Switch

ເປັນຕາກິໂປ່ງແກຣມ

- From user mode to kernel mode

- Interrupts (ເຫັນເຈົ້າຂອງກະທຳ ຜ່ານທ Micro processor)

- Triggered by timer and I/O devices

ກົດຕົວ

- Exceptions (error ex: ຫຍດີຂອງ ອ ມີຮສຫຼຸດ program ຂຶ້ນໄປ)

- Triggered by unexpected program behavior

- Or malicious behavior! ທີ່ໃຫ້ຕົມມີ switch mode

ເປັນຕາກິໂປ່ງແກຣມ

- System calls (aka protected procedure call)

- Request by program for kernel to do some operation on its behalf

ໂຈ່ງຜກລົງນີ້ໃຫຍ່ໃຫ້ kernel ຕົວswitch mode ເພື່ອປຶ້ມມີໃໝ່ເຮັດໄປ

- Only limited # of very carefully coded entry points function ຢັ້ງ

ປະ kernel ຍັກ

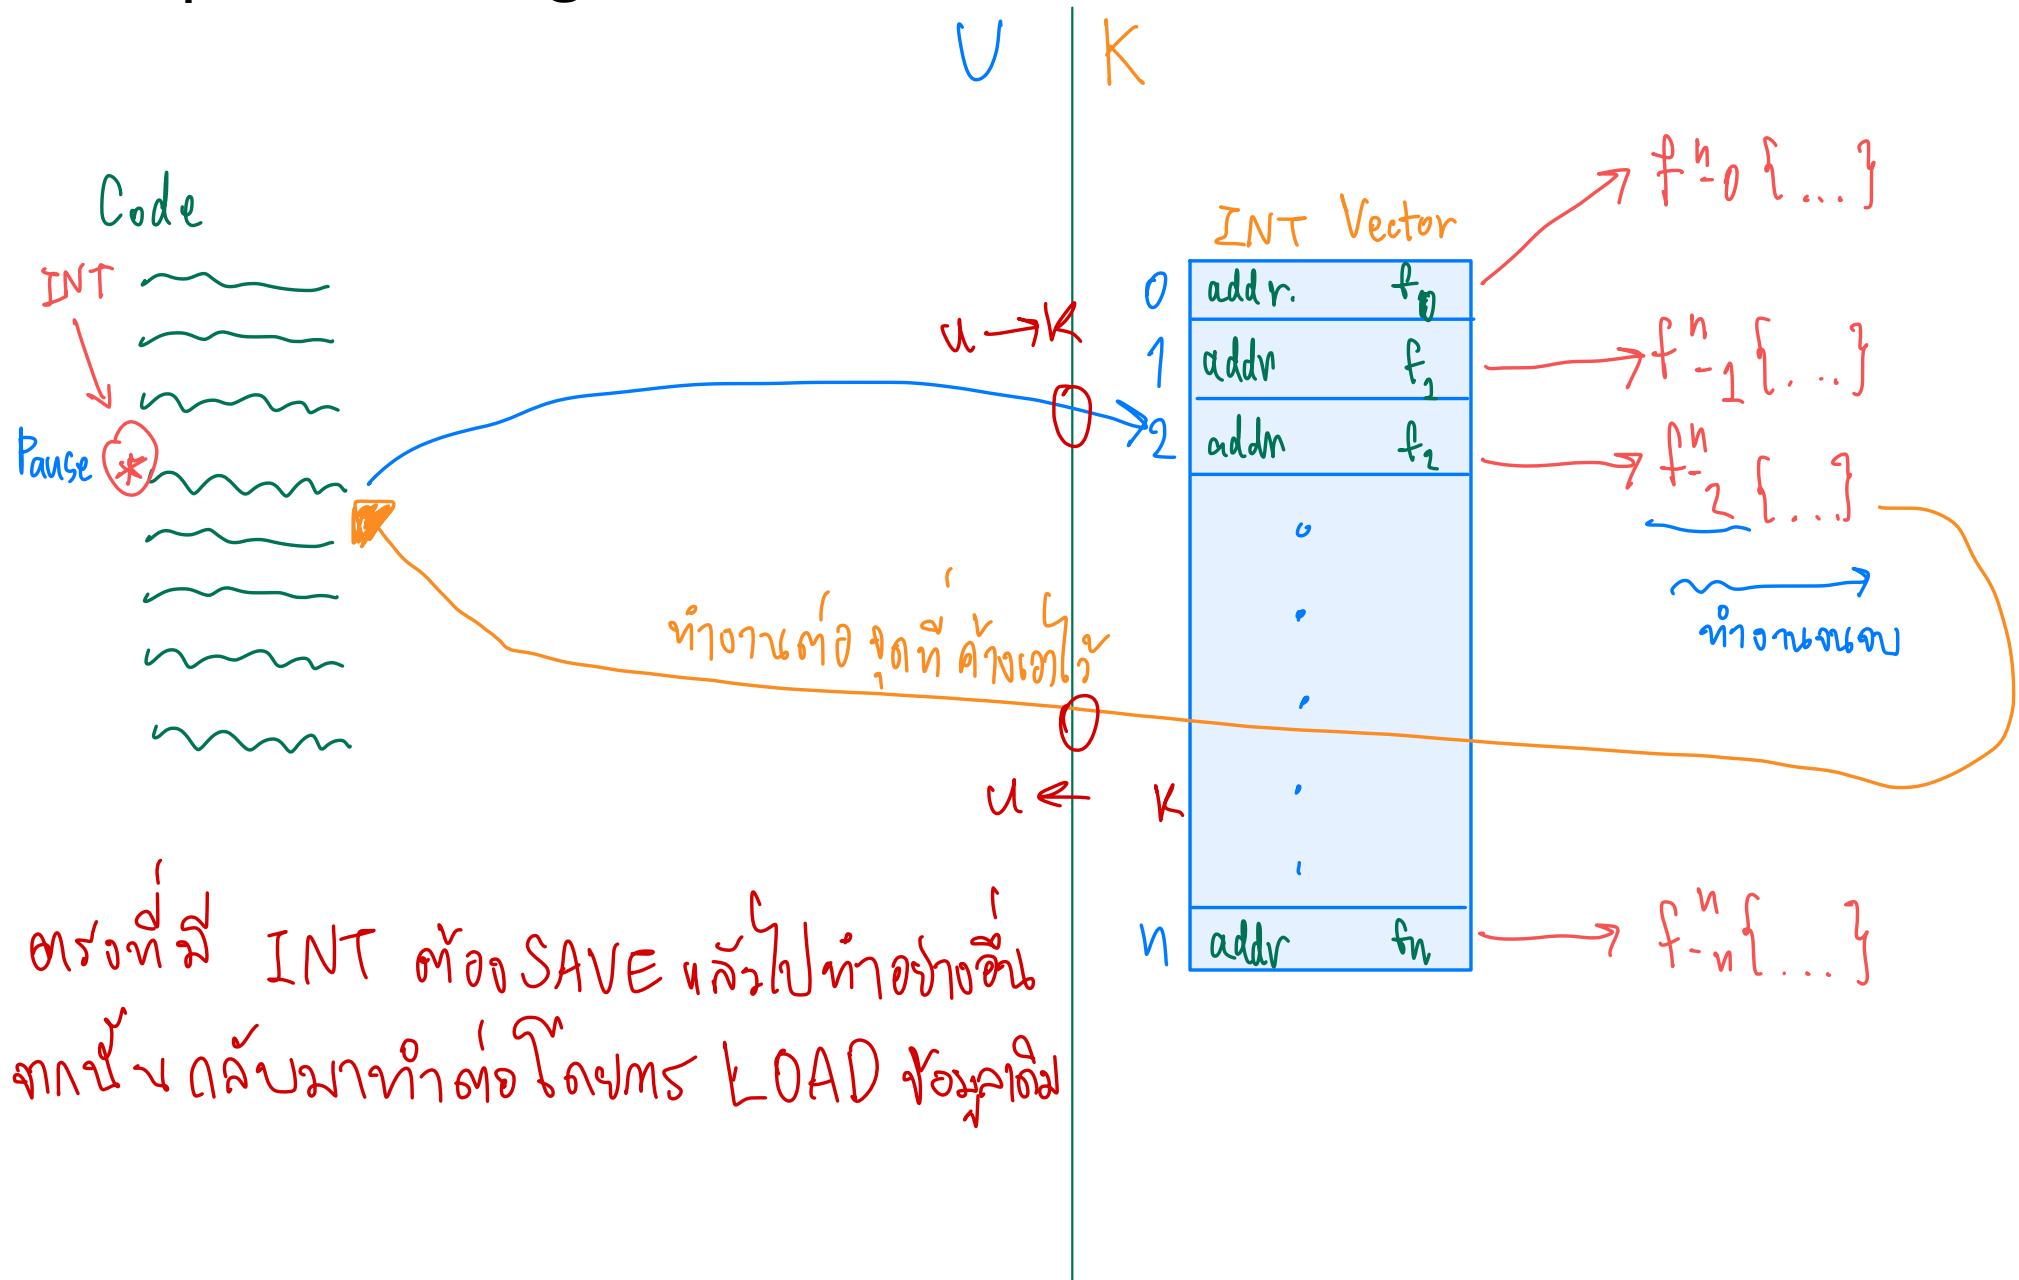
Mode Switch

- From kernel mode to user mode
 - New process/new thread start
 - Jump to first instruction in program/thread
 - Return from interrupt, exception, system call
 - Resume suspended execution
 - Process/thread context switch
 - Resume some other process
 - User-level upcall (UNIX signal)
 - Asynchronous notification to user program
- From kernel mode to user mode
- switch នៅលើកម្រិតការងារ
- Program ការងារ
- switch នៅលើកម្រិតការងារ
- From user mode
- switch នៅលើកម្រិតការងារ
- Kernel mode អាជីវកម្មនៃ function របស់ User → User mode
- ដែលធ្វើឡើងនៅក្នុង Kernel Mode
- ដែលធ្វើឡើងនៅក្នុង Kernel Mode

Activity #1

- ในความเห็นของ นศ การทำ mode switch ควรทำอย่างไรบ้าง เพื่อให้มีความปลอดภัยต่อข้อมูลและเสถียรภาพของระบบ (10 นาที)

Implementing Safe Kernel Mode Transfers



Implementing Safe Kernel Mode Transfers

- Carefully constructed kernel code ^{Save file} packs up the user process state and sets it aside
- Must handle weird/buggy/malicious user state
 - Syscalls with null pointers မျက်ဆုံး
 - Return instruction out of bounds
 - User stack pointer out of bounds
- Should be impossible for buggy or malicious user program to cause the kernel to corrupt itself
- User program should not know interrupt has occurred

(transparency) Programme ၍၂၄၃

ဘာတော်းပေါ်၊ အပြခန္ဓာ ဘုရားမှာ မျှော်လွန်စွာ သာမ်းမှုတော်းပေါ်၊ အလုစုစုပေါ်၊ အချိုင်းအပေါ်၊
၇

Device Interrupts

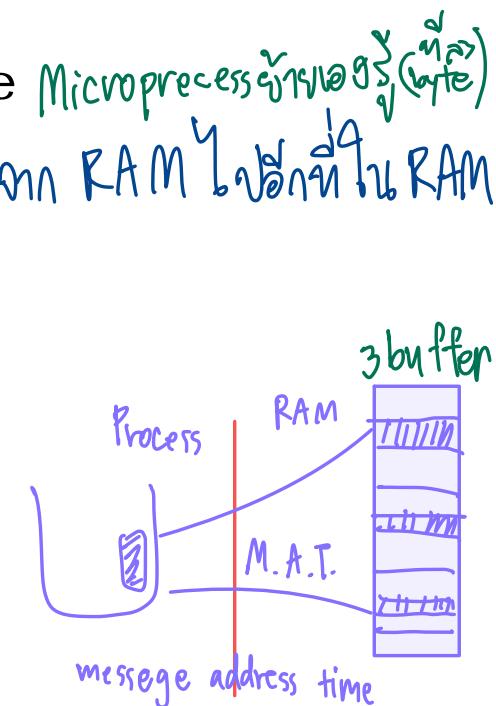
เกี่ยวกับ HW



- OS kernel needs to communicate with physical devices
- Devices operate asynchronously from the CPU
 - Polling: Kernel waits until I/O is done
 - Interrupts: Kernel can do other work in the meantime
- Device access to memory
 - Programmed I/O: CPU reads and writes to device
 - Direct memory access (DMA) by device
 - Buffer descriptor: sequence of DMA's
 - E.g., packet header and packet body
 - Queue of buffer descriptors

Message ที่ไปต่อหน้าเป็นแค่ตัวสินค้า
เพื่อให้ผู้คนสามารถซื้อขายข้อมูล (ค่าใช้จ่าย)
นำไปต่อหน้าร้านค้า ร้านค้า

รับห้อง Buffer | Micro ทำอย่างไรด้วย copy (สร้าง INT ให้)
Micro รับเครื่องไม้ไม่ใช่ copy ดูว่ามีเมื่อไร
ทำอย่างไรต่อไป



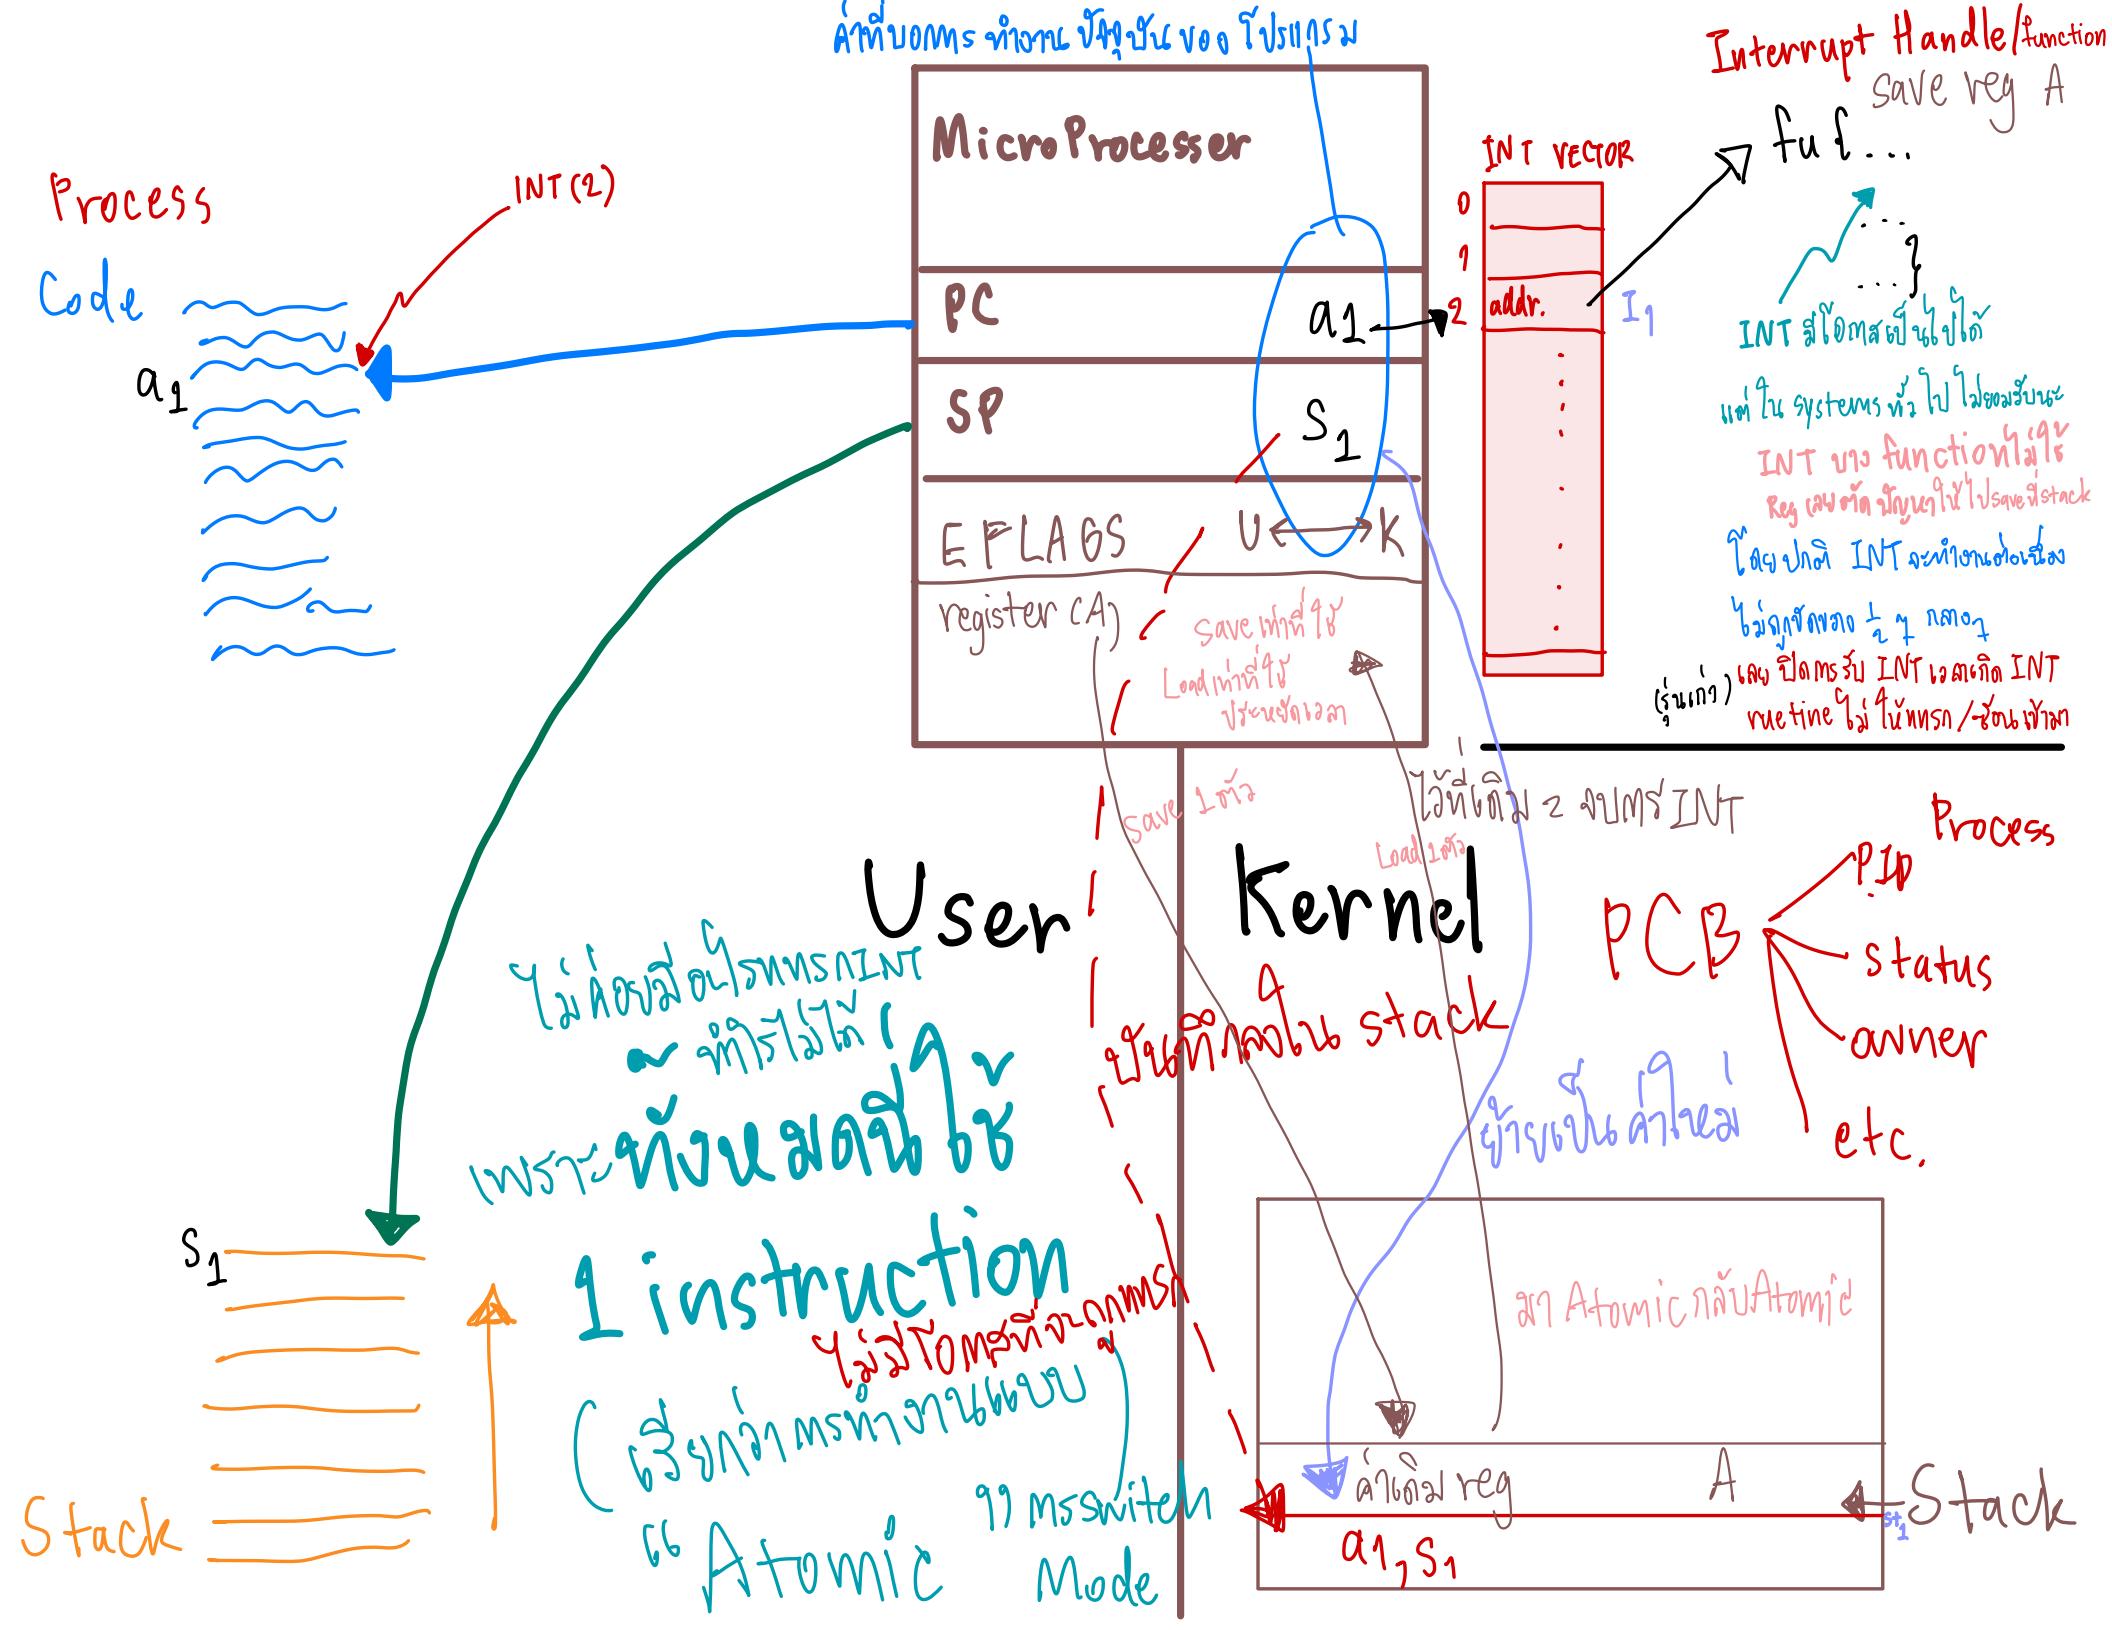
Activity #2

- How do device interrupts work?
 ^{កំណត់ INT នៃវិធាន INT Vector នៅក្នុង}
 ^{ប្រព័ន្ធដែលមានការបញ្ជូនការបន្ថែម}
- Where does the CPU run after an interrupt?
- What stack does it use? ក្នុងការងារ kernel ឬ User
- Is the work the CPU had been doing before the interrupt lost forever? រាយការងារនេះត្រូវបាន Save State ដោយ User
- If not, how does the CPU know how to resume that work? ត្រូវបាន Load State ឡើង

(10 អាទី)

How do we take interrupts safely?

- Interrupt vector
 - Limited number of entry points into kernel
- Kernel interrupt stack
 - Handler works regardless of state of user code
- **Interrupt masking** *in INT*
 - Handler is non-blocking
- Atomic transfer of control
 - “Single instruction”-like to change:
 - Program counter
 - Stack pointer
 - Memory protection
 - Kernel/user mode
- Transparent restartable execution
 - User program does not know interrupt occurred

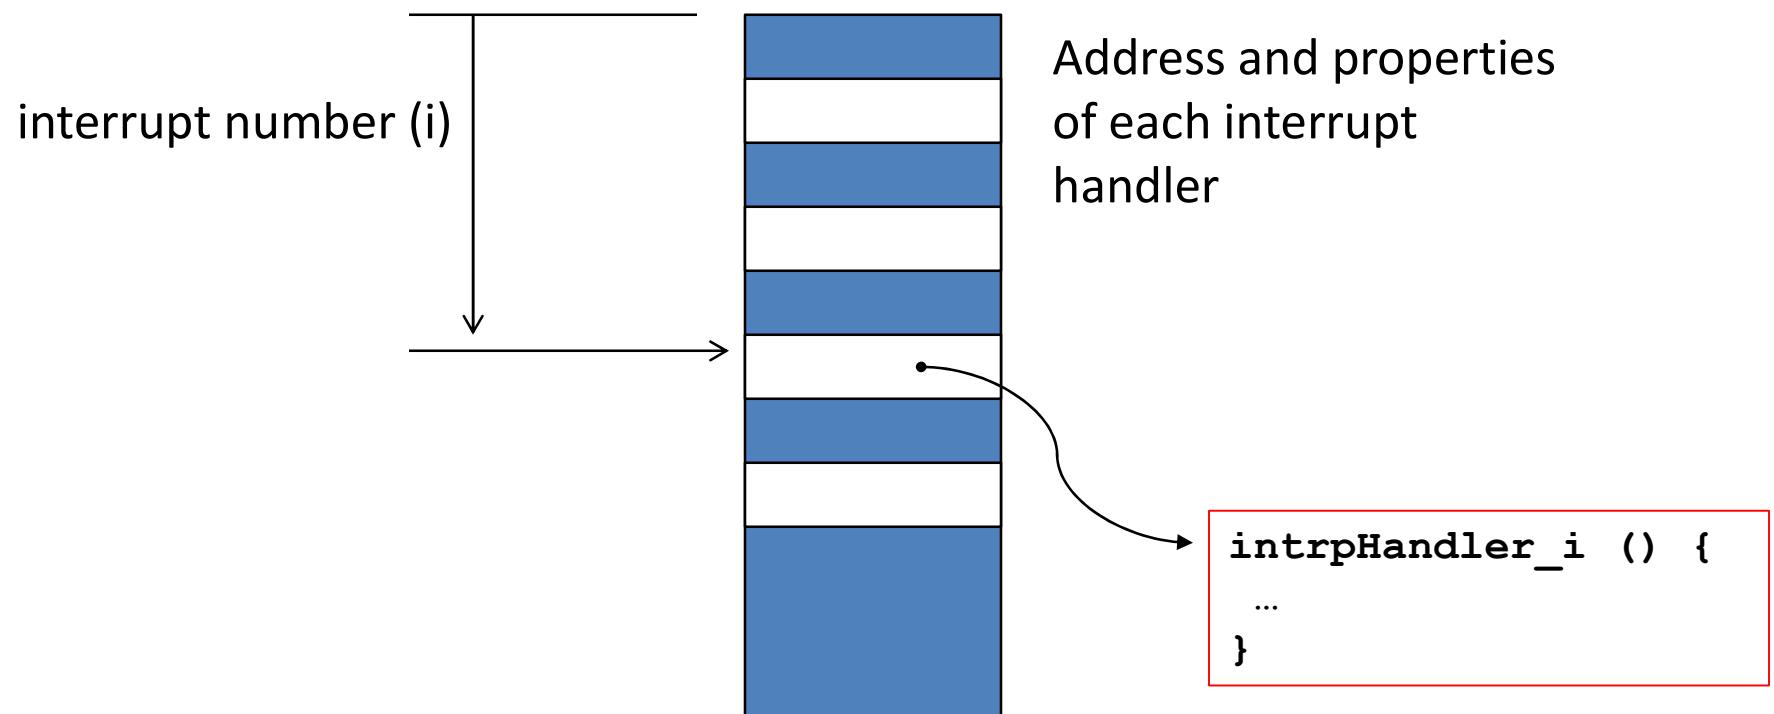


interrupt routine = នៅក្នុងទីតាំងនៃការបញ្ជូនដែលមិនចងក្រោម

- នឹងត្រួវបង្ហាញពី register នៃមេនឹងការ ហើយការណ៍ 1 register (នឹងរួច) ការសំនើនៃការងារនេះ save register A ឬនៅក្នុងក្រុងការណ៍ និងការងារនេះ save register stack

Where do mode transfers go?

- Solution: ***Interrupt Vector***



The Kernel Stack

เฝ้า Kernel

- Interrupt handlers want a stack

A. เสื่อมรุกพาพิจาระบบ
ระบบช่วยเหลือ

- System call handlers want a stack

- Can't just use the user stack [why?]

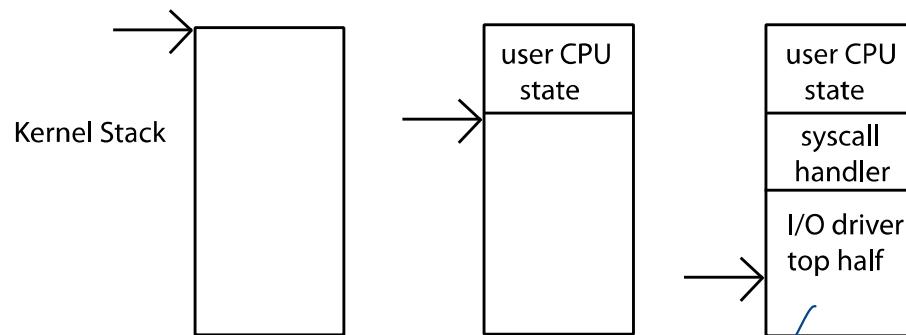
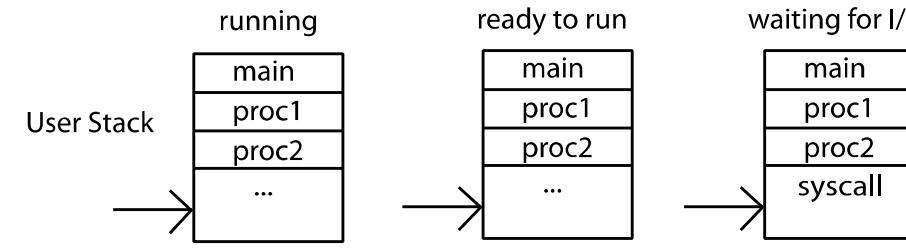
สามารถ detect ms INT ได้

ใน 1 process ล้วน เดียว thread แค่ ก็มี INT thread

อีกอันก็ทำงานได้ แต่เรา INT ไม่ detect ค่าที่ไม่ใช่
ผู้ใช้ ต้องที่อยู่ตัว พร้อมปีน stack หาด้วย User แล้วเดา
ถูกๆ ก็จะ แก้ไขได้

The Kernel Stack

- Solution: two-stack model
 - Each OS thread has kernel stack (located in kernel memory) plus user stack (located in user memory)
- Place to save user registers during interrupt



Interrupt Stack

- Per-processor, located in kernel (not user) memory
 - Usually a process/thread has both: kernel and user stack
- Why can't the interrupt handler run on the stack of the interrupted user process?

Case Study: x86 Interrupt

Atomic
Instruction
Instruction

- Save current stack pointer
- Save current program counter
- Save current processor status word (condition codes)
- Switch to kernel stack; put SP, PC, PSW on stack
- Switch to kernel mode
- Vector through interrupt table
- Interrupt handler saves registers it might clobber

Save reg $\times 9$

Save $\times 9$ load $\times 9$ $\times 9$

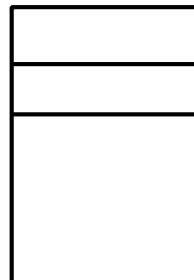
Before Interrupt

User-level
Process

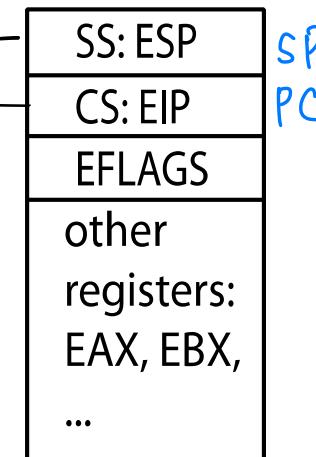
code:

```
foo () {  
    while(...) {  
        x = x+1;  
        y = y-2;  
    }  
}
```

stack:



Registers

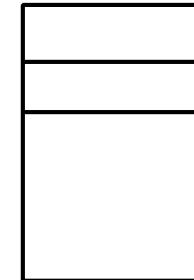


Kernel

code:

```
handler() {  
    pusha  
    ...  
}
```

Exception
Stack



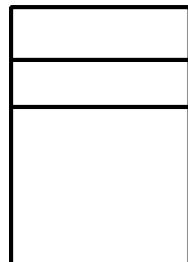
During Interrupt

User-level
Process

code:

```
foo () {  
    while(...) {  
        x = x+1;  
        y = y-2;  
    }  
}
```

stack:



Registers

SS: ESP
CS: EIP
EFLAGS
other
registers:
EAX, EBX,
...

Kernel

code:

```
handler() {  
    pusha  
    ...  
}
```

Exception
Stack

SS
ESP
EFLAGS
CS
EIP
error

After Interrupt

User-level
Process

code:

```
foo () {  
    while(...) {  
        x = x+1;  
        y = y-2;  
    }  
}
```

stack:



Registers

SS: ESP
CS: EIP
EFLAGS
other registers: EAX, EBX, ...

Kernel

code:

```
After INT  
handler() {  
    pusha  
    ...  
}
```

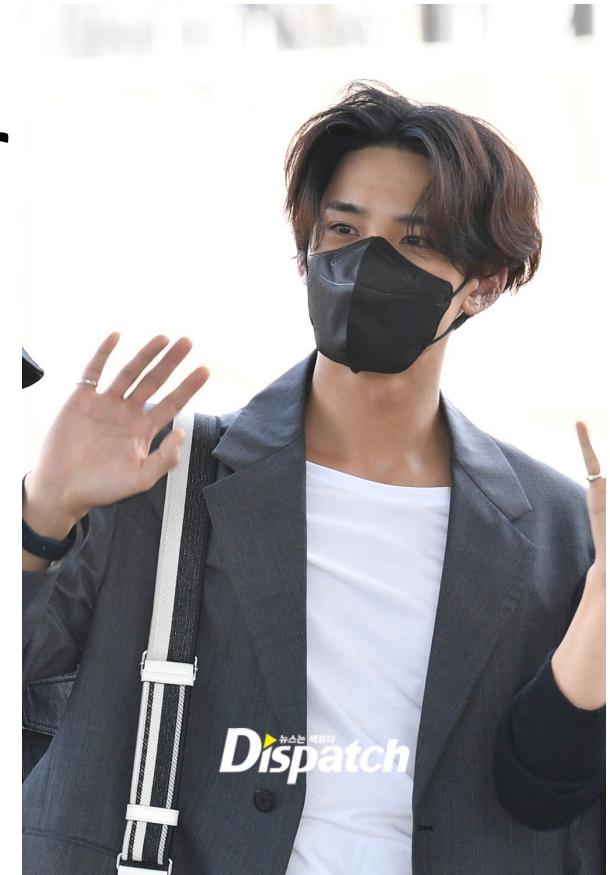
After save a register save a register

Exception
Stack

SS
ESP
EFLAGS
CS
EIP
error
Other Regs.

At end of handler

- Handler restores saved registers
- Atomically return to interrupted process/thread
 - Restore program counter
 - Restore program stack
 - Restore processor status word/condition codes
 - Switch to user mode



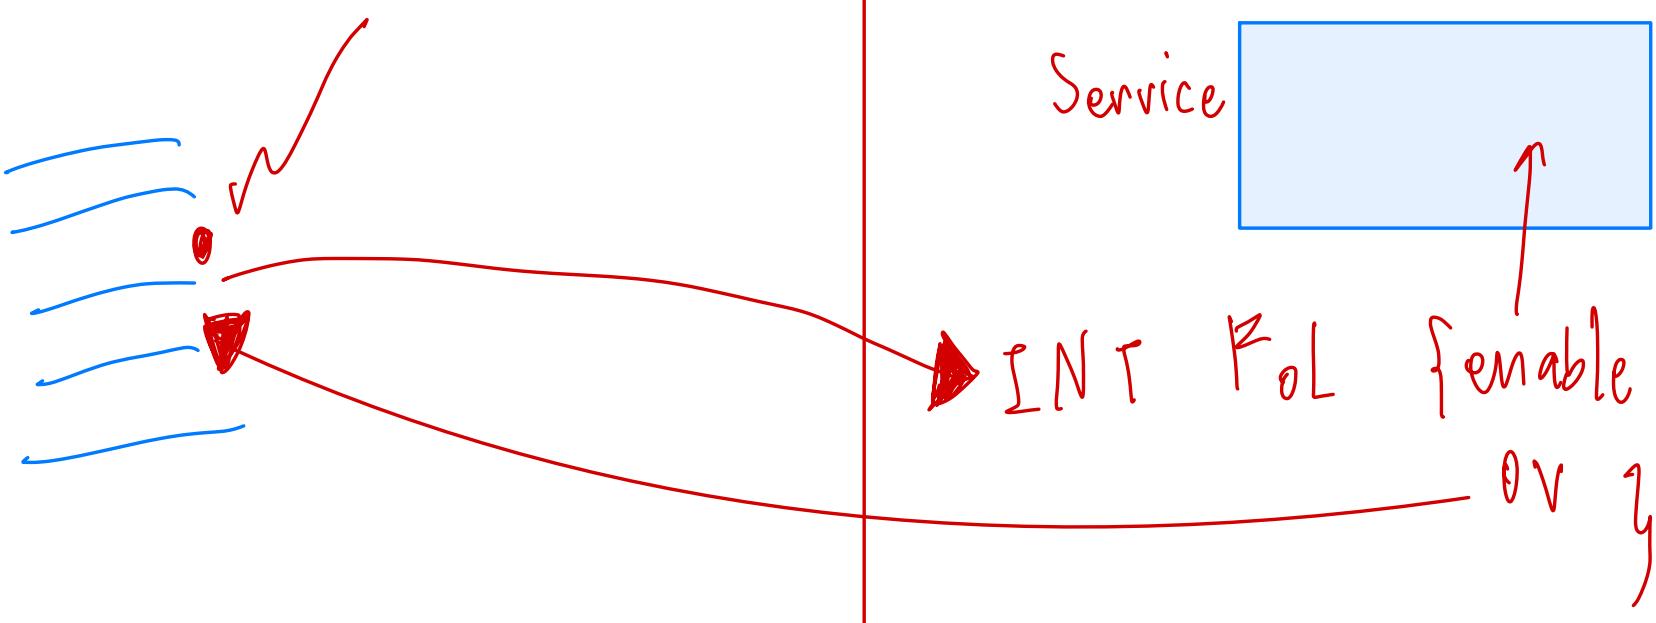
Interrupt Masking

ក្រោមឯង ទាំងនៅ interrupt នេះ ត្រូវរាយការណ៍របៀប INT

- Interrupt handler runs with interrupts off
 - Re-enabled when interrupt completes
- OS kernel can also turn interrupts off
 - Eg., when determining the next process/thread to run
 - On x86
 - CLI: disable interrupts
 - STI: enable interrupts
 - Only applies to the current CPU (on a multicore)
- We'll need this to implement synchronization in chapter 5

Hardware support: Interrupt Control

- Interrupt processing not visible to the user process:
 - Occurs ^{เกิด} between instructions, restarted transparently
 - No change to process state $\text{SAVE} \rightarrow \text{LOAD}$ กลับมา
 - What can be observed even with perfect interrupt processing? คำตอบ รีดิจ์ สังข์กันโน่ โดยทั่วไป ปุ่มหนึ่งกด 100 ครั้ง ก็จะเห็นว่ามันยังคงเด้งอยู่ (ยกเว้น INT ก็จะหายไป)
แล้วเที่ยบผลกับที่เราต้องการว่ามันผลลัพธ์ตรง (ถ้า INT ก็จะรีบูตกลับไป)
- Interrupt Handler invoked with interrupts ‘disabled’
 - Re-enabled upon completion
 - Non-blocking (run to completion, no waits)
 - Pack up in a queue and pass off to an OS thread for hard work
 - wake up an existing OS thread



Hardware support: Interrupt Control

- OS kernel may enable/disable interrupts
 - On x86: CLI (disable interrupts), STI (enable)
 - Atomic section when select next process/thread to run
 - Atomic return from interrupt or syscall
- HW may have multiple levels of interrupts (ក្នុងប្រព័ន្ធ)
 - Mask off (disable) certain interrupts, eg., lower priority
INT ត្រូវបាន ស្វែងស្ថិក
 - Certain Non-Maskable-Interrupts (NMI)
 - e.g., kernel segmentation fault
មីនាក់តិចទៅរួមបញ្ជី ដើម្បីបង្ហាញការកំណត់ពាណិជ្ជកម្ម
 - Also: Power about to fail!
(ធម៌មានបញ្ហាលើ ពេលវេលាដឹងថា មីនា SAVE ឡើង)

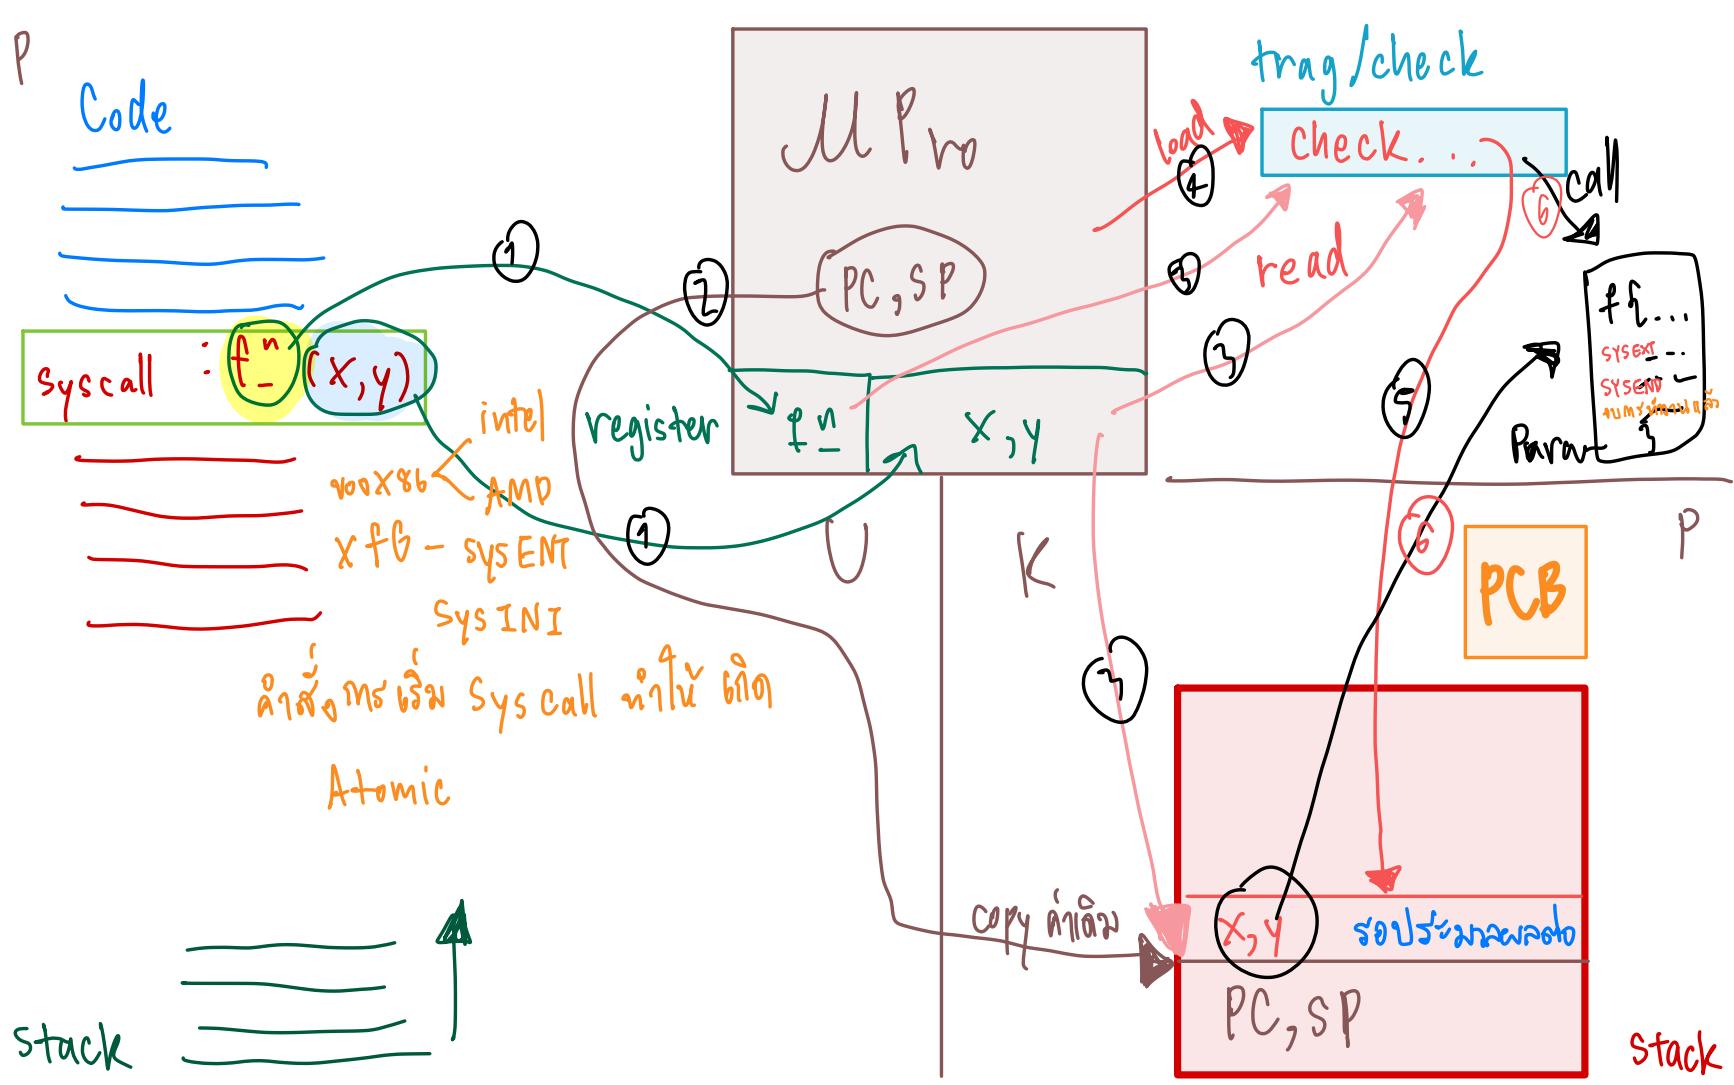
INT Level ខ្ពស់

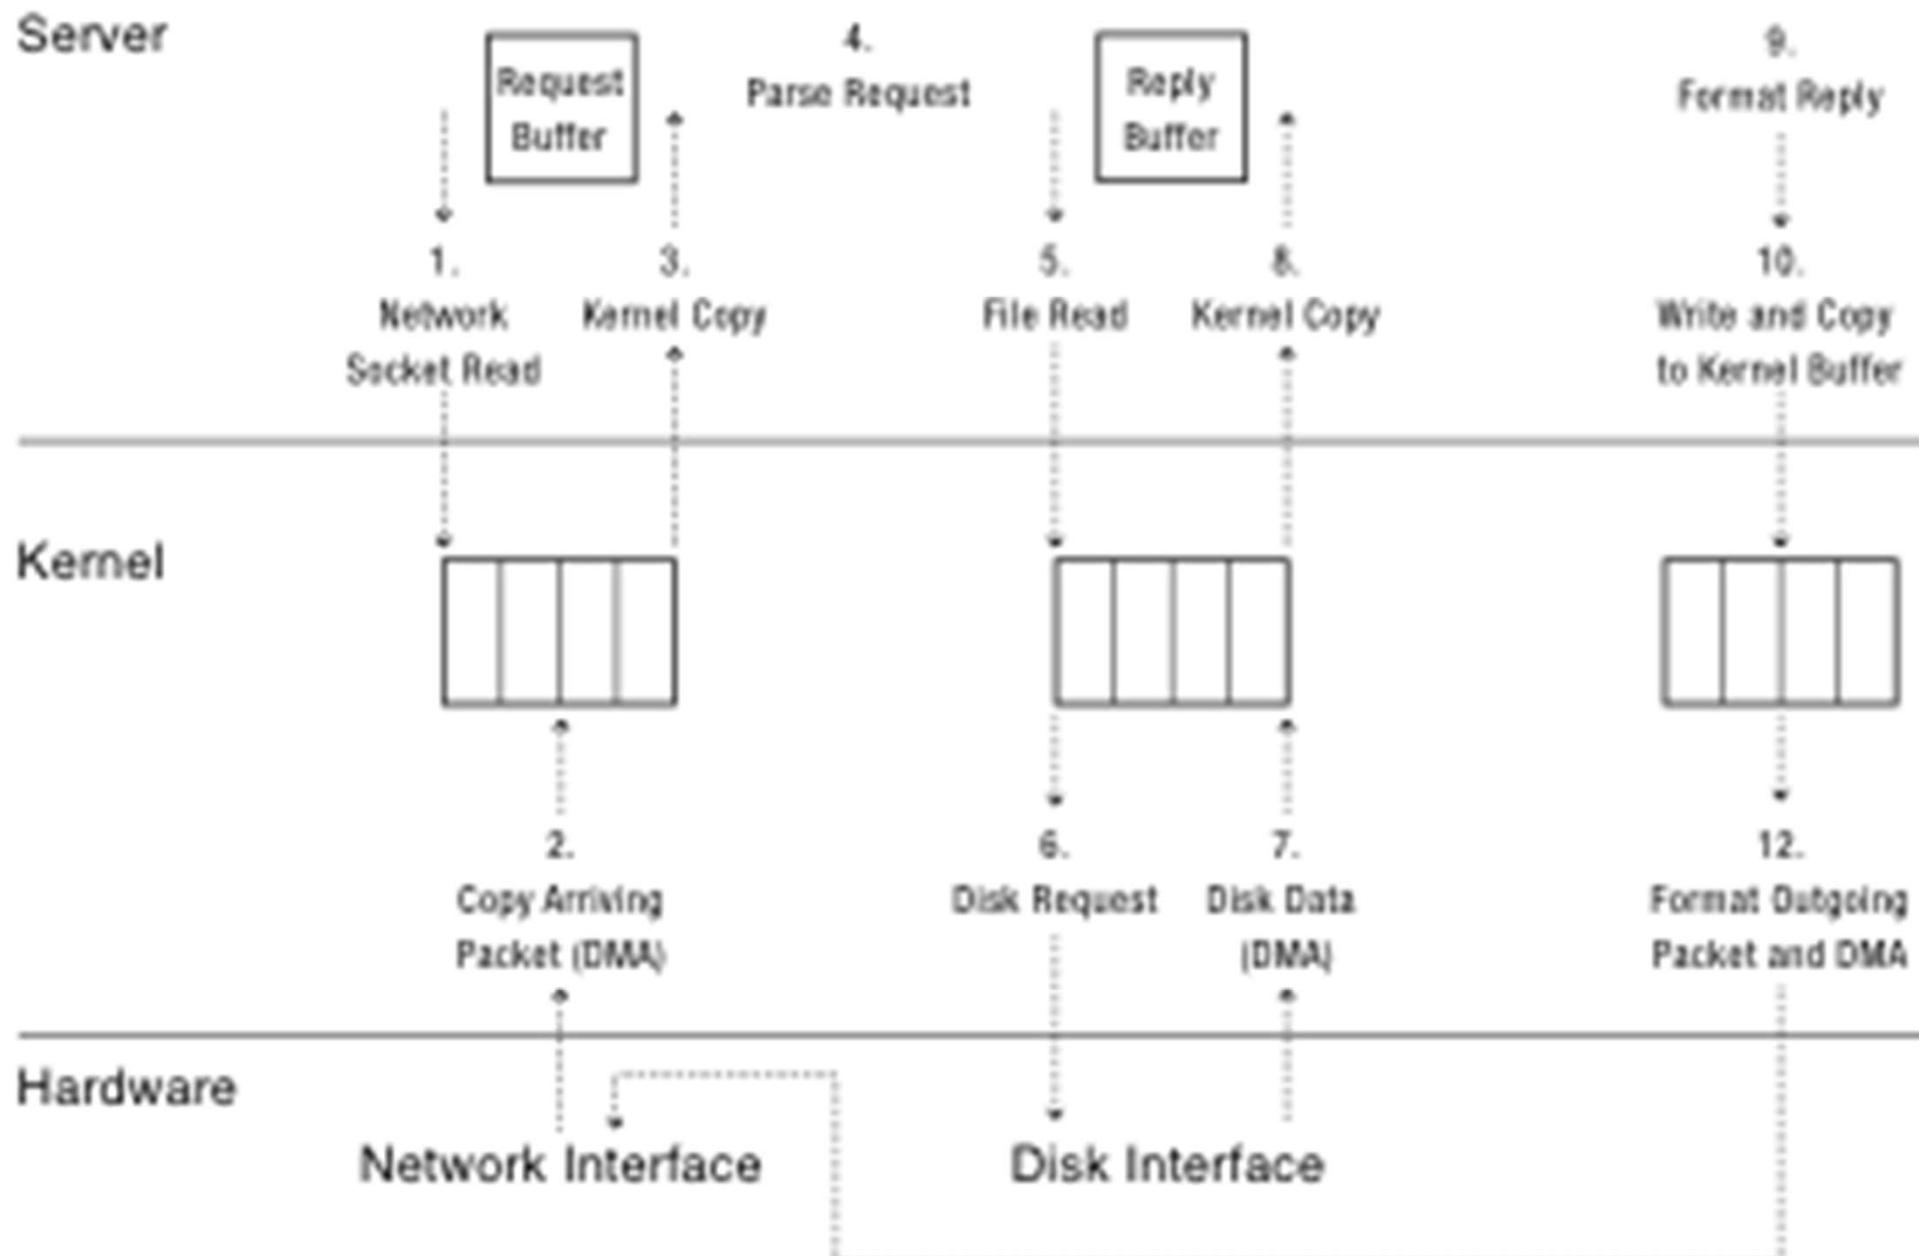
សម្រាប់ INT ខ្ពស់

Level ពីរគ្មាន

Kernel System Call Handler

- Vector through well-defined syscall entry points!
 - Table mapping system call number to handler
- Locate arguments
 - In registers or on user (!) stack
- Copy arguments
 - From user memory into kernel memory – carefully checking locations!
 - Protect kernel from malicious code evading checks
- Validate arguments
 - Protect kernel from errors in user code
- Copy results back
 - Into user memory – carefully checking locations!





Virtual Addr = Logic Addr.

Today: Four Fundamental OS Concepts

- **Thread: Execution Context**
 - Program Counter, Registers, Execution Flags, Stack
- **Address space (with translation)**
 - Program's view of memory is distinct from physical machine
- **Process: an instance of a running program**
 - Address Space + One or more Threads
- **Dual mode operation / Protection**
 - Only the “system” can access certain resources
 - Combined with translation, isolates programs from each other

