


P $\xrightarrow{\text{fork}}$ C

segment	B	Bo	Access
Code	B ₀	A ₀	R
data	B ₁	A ₁	R/W
heap	B ₂	A ₂	R/W
stack	B ₃	A ₃	R/W

segment share
↑
no copy
set Robt
R

segment	B	Bo	Access
Code	B ₀	A ₀	
data	B ₁	A ₁	
heap	B ₂	A ₂	
stack	B ₃	A ₃	



Virtual Memory

COPY ON

WRITE


- ① ใช้หน่วยความจำ RAM
 - ② พอหน่วยความจำเต็ม (ready to write)
- ถ้าไม่มีการ write (เขียน) ก็ยัง share อยู่



Fill on Reference

ถ้าใช้วิธีนี้เรื่องการใช้ function ก็ไม่มีปัญหาแล้ว memory mapped file เป็น block เข้ากับ frame ของ RAM



Logical



load
main() {
ไปรันพร้อมทำงานเร็วขึ้น
ใช้ RAM ของระบบ/โปรแกรม


zero copy I/O
บอกรุ่น (page, frame)
ที่ไม่ใช่ copy เก็บ
(background)

Storage (hardisk)




Definitions

- Cache
 - Copy of data that is faster to access than the original
 - Hit: if cache has copy
 - Miss: if cache does not have copy
- Cache block แบ่งเป็น block
 - Unit of cache storage (multiple memory locations)
- Temporal locality ทรงใช้บ่อยๆ ในลักษณะ loop
 - Programs tend to reference the same memory locations multiple times ลักษณะการนำเอาค่าของ array มาใช้ซ้ำๆ ในโปรแกรม ยกเว้นเมื่อ Call function
 - Example: instructions in a loop
- Spatial locality บริเวณเดียวกัน
 - Programs tend to reference nearby locations
 - Example: data in a loop กับ Array



Memory mapped file



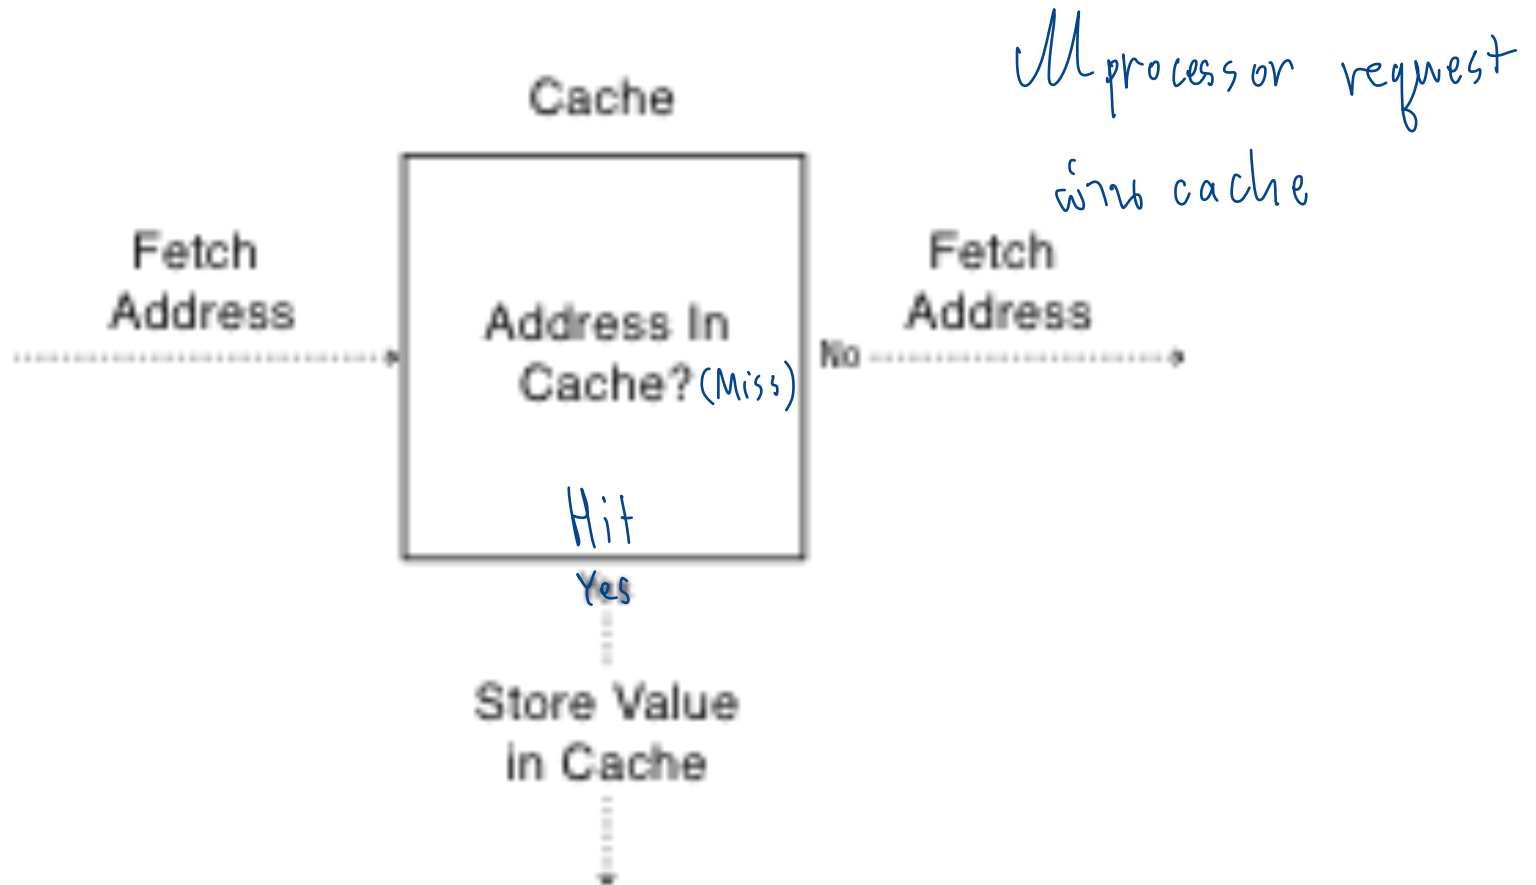
	frame	access
Row/Entry 1	1	R
77		invalid

invalid 77 slot for RAM only

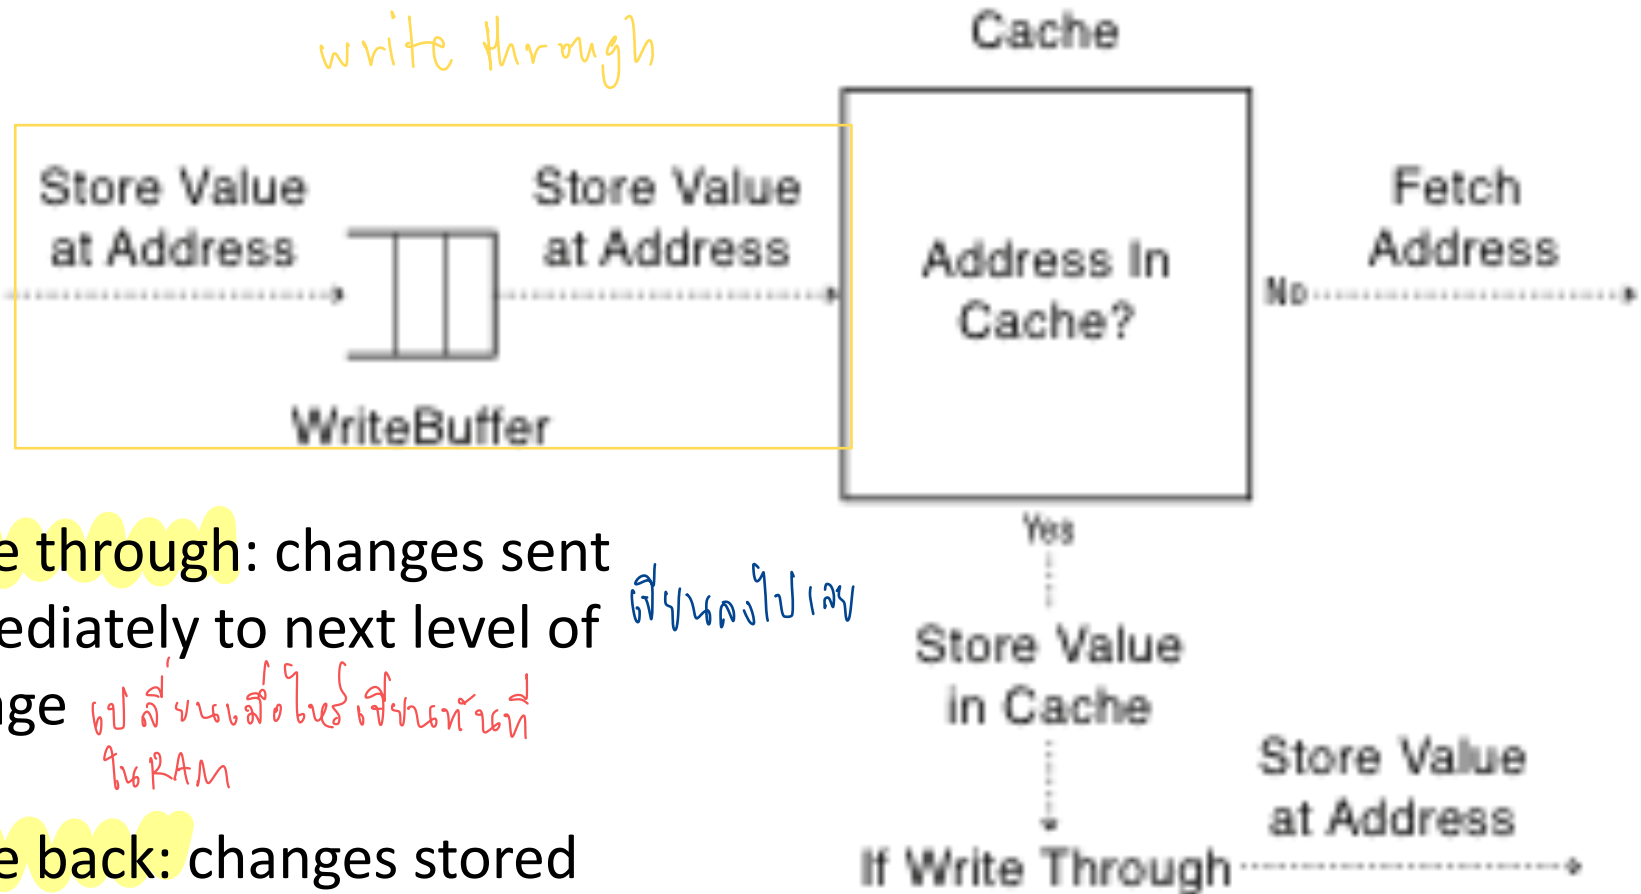
Storage

	frame	access
Row/Entry 1	1	R
Row/Entry 3	3	R/W

Cache Concept (Read)



Cache Concept (Write)



Write through: changes sent immediately to next level of storage
เขียนลงไปแล้ว
เปลี่ยนเนื้อหาที่เขียนทันทีใน RAM


Write back: changes stored in cache until cache block is replaced
เขียนใน cache พอเต็มค่อยทำ I/O กับ RAM ในขั้นต่อไป
เอาออกไปเก็บใน RAM (update ใน RAM)

Memory Hierarchy


Cache	Hit Cost	Size
1st level cache/first level TLB	1 ns	64 KB
2nd level cache/second level TLB	4 ns	256 KB
3rd level cache	12 ns	2 MB
Memory (DRAM)	100 ns	10 GB
Data center memory (DRAM)	100 μ s	100 TB
Local non-volatile memory	100 μ s	100 GB
Local disk	10 ms	1 TB
Data center disk	10 ms	100 PB
Remote data center disk	200 ms	1 XB

i7 has 8MB as shared 3rd level cache; 2nd level cache is per-core

Cache hit rate



Example of cache hit rate over time




Main Points

- Can we provide the illusion of near infinite memory in limited physical memory?
 - Demand-paged virtual memory
 - Memory-mapped files
- How do we choose which page to replace?
 - FIFO, MIN, LRU, LFU, Clock
- What types of workloads does caching work for, and how well?
 - Spatial/temporal locality vs. Zipf workloads


Hardware address translation is a power tool

- Kernel trap on read/write to selected addresses
 - Copy on write *ถ้ามี read/write จะ copy ออกไป free space*
 - Fill on reference *translation น่าใช้*
 - Demand paged virtual memory
 - Memory mapped files
 - Modified bit emulation
 - Use bit emulation
- ใน RAM LOAD ข้อมูลที่ Programerกำหนดไว้*

Demand Paging (Before)



Demand Paging (After)



Demand Paging

1. TLB miss
2. Page table walk
3. Page fault (page invalid in page table)
4. Trap to kernel
5. Convert virtual address to file + offset
6. Allocate page frame
 - Evict page if needed
7. Initiate disk block read into page frame
8. Disk interrupt when DMA complete
9. Mark page as valid
10. Resume process at faulting instruction
11. TLB miss
12. Page table walk to fetch translation
13. Execute instruction

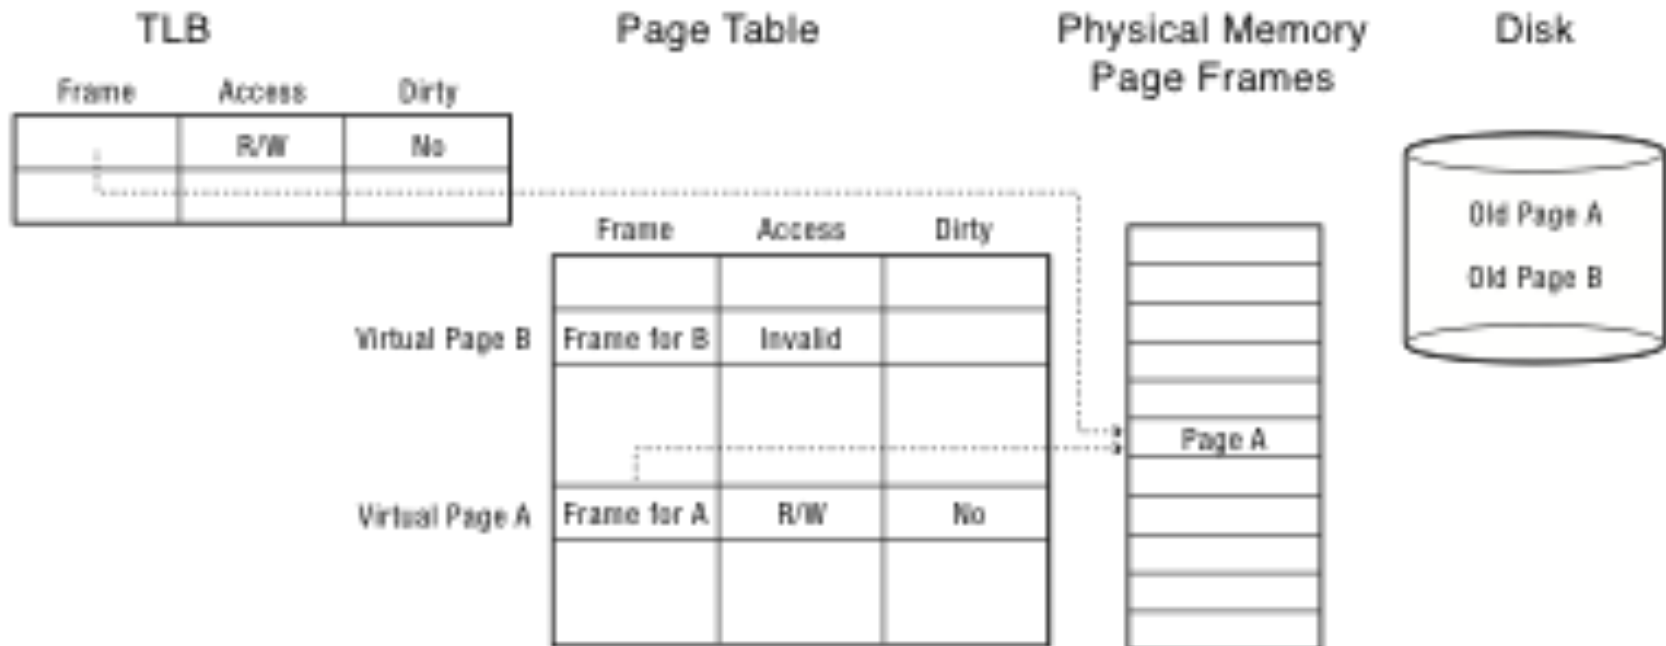
Allocating a Page Frame

- Select old page to evict
- Find all page table entries that refer to old page
 - If page frame is shared
- Set each page table entry to invalid
- Remove any TLB entries
 - Copies of now invalid page table entry
- Write changes on page back to disk, if necessary


How do we know if page has been modified?

- Every page table entry has some bookkeeping
 - Has page been modified?
 - Set by hardware on store instruction
 - In both TLB and page table entry
 - Has page been recently used?
 - Set by hardware on in page table entry on every TLB miss
- Bookkeeping bits can be reset by the OS kernel
 - When changes to page are flushed to disk
 - To track whether page is recently used

Keeping Track of Page Modifications (Before)



Keeping Track of Page Modifications (After)



Virtual or Physical Dirty/Use Bits

- Most machines keep dirty/use bits in the page table entry
- Physical page is
 - Modified if *any* page table entry that points to it is modified
 - Recently used if *any* page table entry that points to it is recently used
- On MIPS, simpler to keep dirty/use bits in the core map
 - Core map: map of physical page frames

Models for Application File I/O

- Explicit read/write system calls
 - Data copied to user process using system call
 - Application operates on data
 - Data copied back to kernel using system call
- Memory-mapped files
 - Open file as a memory segment
 - Program uses load/store instructions on segment memory, implicitly operating on the file
 - Page fault if portion of file is not yet in memory
 - Kernel brings missing blocks into memory, restarts process

Advantages to Memory-mapped Files

- Programming simplicity, esp for large files
 - Operate directly on file, instead of copy in/copy out
- Zero-copy I/O
 - Data brought from disk directly into page frame
- Pipelining
 - Process can start working before all the pages are populated
- Interprocess communication
 - Shared memory segment vs. temporary file

Cache Replacement Policy

- On a cache miss, how do we choose which entry to replace?
 - Assuming the new entry is more likely to be used in the near future
 - In direct mapped caches, not an issue!
- Policy goal: reduce cache misses
 - Improve expected case performance
 - Also: reduce likelihood of very poor performance

A Simple Policy

- Random?
 - Replace a random entry
- FIFO?
 - Replace the entry that has been in the cache the longest time
 - What could go wrong?

FIFO in Action

FIFO

Reference	A	B	C	D	E	A	B	C	D	E	A	B	C	D	E
1	A				E				D				C		
2		B				A				E				D	
3			C				B				A				E
4				D				C				B			

Worst case for FIFO is if program strides through memory that is larger than the cache

MIN, LRU, LFU

- MIN
 - Replace the cache entry that will not be used for the longest time into the future
 - Optimality proof based on exchange: if evict an entry used sooner, that will trigger an earlier cache miss
- Least Recently Used (LRU)
 - Replace the cache entry that has not been used for the longest time in the past
 - Approximation of MIN
- Least Frequently Used (LFU)
 - Replace the cache entry used the least often (in the recent past)

LRU/MIN for Sequential Scan

LRU															
Reference	A	B	C	D	E	A	B	C	D	E	A	B	C	D	E
1	A				E				D				C		
2		B				A				E				D	
3			C				B				A				E
4				D				C				B			
MIN															
1	A					+					+			+	
2		B					+					+	C		
3			C					+	D					+	
4				D	E					+					+

LRU															
Reference	A	B	A	C	B	D	A	D	E	D	A	E	B	A	C
1	A		+				+				+			+	
2		B			+								+		
3				C					E			+			
4						D		+		+					C
FIFO															
1	A		+				+		E						
2		B			+						A			+	
3				C								+	B		
4						D		+		+					C
MIN															
1	A		+				+				+			+	
2		B			+								+		C
3				C					E			+			
4						D		+		+					

Belady's Anomaly

FIFO (3 slots)												
Reference	A	B	C	D	A	B	E	A	B	C	D	E
1	A			D			E					+
2		B			A			+		C		
3			C			B			+		D	
FIFO (4 slots)												
1	A				+		E				D	
2		B				+		A				E
3			C						B			
4				D						C		

Recap

- MIN is optimal
 - replace the page or cache entry that will be used farthest into the future
- LRU is an approximation of MIN
 - For programs that exhibit spatial and temporal locality